

Week09 development report - Hongle Yang

CS30-1 personalised AI learning assistant | 6 October 2026

Status: Current source and recorded development results.

Accountable domain: Real textbook provenance, extraction and source structure. Original references: DAT-10, DAT-12. Current trace focus: REQ-W09-05-01 to 05-04; 03-01 and 03-02. Named accountability does not certify personal execution.

Domain methods and contracts

The source baseline consists of the actual four OpenStax books and the retained 10,594-vector release. Source dates, licences, original file hashes, version identity, text spans, page locators, coordinates, units and indices must survive reprocessing. Authored fixtures can test mechanics but cannot replace construction from the real books. This drafting task performed no current database count or corpus-content audit.

Structured extraction separates successful parsing from semantic review. A table needs correct cells and headers, a formula needs its glyphs, indices and units, and the original page geometry and reading order must remain inspectable. A retained full 185-MB isolated parse is historical evidence for that version, with bounded exception statistics; it is not automatic approval of every extracted region. Concept and prerequisite relationships require attributed review before publication.

The case inventory supplies deterministic input truths. G01 has no supplied candidate and leaves a first inverse-choice judgment pending. S11 has a supplied candidate and fixed gas amount and temperature, with 100 kPa, 2 L and 4 L among the explicit givens. Evidence sufficiency and the difference between an explicit given and an inferred consequence must remain visible to the checker and reviewer.

Actual findings and error analysis

The earlier S11 D07 response described pressure change as part of the givens, although a final pressure was not explicitly supplied; that attribution concern remains alongside its original PASS flags. D09 then produced mixed reviews for the same source version. Exact quote containment proves that a quote exists, but does not certify every compound claim or the pedagogical relevance of displaying it. A later result must not rewrite an earlier source or judgment.

The final package needs the documented real-source closure or an honest dependency/reconstruction requirement. The older 951 source-review ZIPs omit large source/runtime materials and cannot be called the current complete installation. Original PDFs, earlier vector releases and relation-review history must remain preserved rather than being replaced by convenience extracts.

Source and verification boundary

The current release contains 972 canonical source inputs, digest ``c9a35bfeee65acbf2e7dd2c8fcdc1df86ef89f71af5781873d3c7581145389c9``. Its complete eight-stage free software gate passed: 4,107 base Python cases passed, 102 additional subchecks were reported, and one original optional case was skipped. All 483 integration cases passed. The unit suite had 3,624 passes; its 19 current question-boundary cases are included in that count. Frontend verification passed

204 cases in 22 files. Every stage exited successfully, with zero recorded failures or errors. These observations are software checks, not formal teaching-quality or performance estimates.

The actual D09 native source is

``d6888ca8771646fd7c59bf9869d63dc80b620bf583c8863ff11aceda7845ff9b``. Its earlier full gate had seven passing stages, 4,082 base Python passes, 25 integration failures, one optional skip and 102 additional subchecks; its unit suite passed 3,624 cases, integration passed 458/483, and frontend passed 204. This failed gate is preserved. The test-only 6748 successor passed 28 focused cases. The final formatted cleanup test retained the tested syntax tree and passed Ruff; an additive handover record completed the c9 source capture.

The exact production compatibility bridge

``380a5d71e12fc2ba5c030efeafeab936567d843916c63e57bbbf7ccb4425446`` verifies identical bytes for 639 production Python, prompt, frontend and configuration files between the D09 source and current c9. The whole-source hashes differ. Native D09 results remain attached to their actual source; the bridge does not supply a new model run or prove improved semantic quality.

The current public operational payload has 1,059 verified files and digest

``afca072d76faffdf7d0ba599c6d43b8bfb52ca82ef90396939262de74cfb006``, including all 83 public resource members. Dependency installation uses locked online dependencies and host prerequisites. Twelve opaque research fixtures are excluded from the public payload, so it does not reconstruct every canonical test input. Production and operational files remained unchanged during the installed checks.

Installed observations and preserved chronology

The original fresh CPU installation remains FAILED. Its unchanged launcher exited successfully, dependency checks passed, four roles became ready, and the official-corpus query observed four documents, one active entry, zero model credentials and zero model configurations.

Browser/workspace checks had not run when it stopped. Its recorded ValueError cause is unproven. The denied original credential remains untouched. Later continuations did not read or stat it, alter ACLs, or try elevated or alternate access; they used new supported anonymous actors. Its contents and size were not established, and no product password defect is inferred. The default bootstrap-credential journey remains unverified.

Continuation 01 failed before actor creation because its fixture used a reserved email domain.

Continuation 02 created supported anonymous actors and started four roles, then failed clean-environment Edge resolution before browser assertions. Continuation 03 passed all five connected browser cases in 7.3 seconds but its workspace failed at an exact nested-label locator after six completed checks. Continuation 04 preserved that browser result without rerunning it; its workspace stopped at a saved-and-due-now notice literal mismatch. All five original FAILED lifecycle receipts remain unchanged.

Continuation 05 actually passed the declared installed connected checks. Its entire workspace passed 11 assertions, recorded 19 screenshots and reported zero page_errors entries. The five browser cases are reused explicitly from continuation 03: three original specifications were unchanged, and one security specification changed only an anonymous admin-email literal while preserving every assertion. The supported anonymous actors and four direct role entry points use the same installed c9 source. The current workspace verifies PDF/source identity, reading selection and bookmark, practice rules and pending short/step assessments, review scheduling, notes and both export formats, practice-derived memory provenance, sign-out isolation, manual draft validation/publication, and a read-only admin visual candidate. These checks do not certify all eight modules or live teaching quality.

Review of all 19 workspace screenshots observed a `CONNECTION_ERROR` in the goal-desktop session sidebar, while goal saving/readback and the 11 runtime assertions passed. Its fetch-rejection mapping is known, but requestfailed telemetry was not captured and the cause remains unproven; recorded session/API responses were successful, so a persistent server outage is not established. Other capture limits remain: transient Loading in the desktop admin image, partially occluded reader-mobile top controls, and source-overlay coordinate alignment that pixels alone cannot prove. An empty `page_errors` list does not establish error-free screens or complete visual acceptance.

The original recovery attempt 10 remains FAILED. Backup and restore ran, but the strict original-source versus restored full-65-table aggregate differed at `retrieval_epochs.token` and the API restart phase did not start. Read-only diagnosis found the intended `document_versions` restore trigger invalidated this cache epoch with a fresh UUID; the other 64 tables and four original source files matched exactly. This comparison error does not relabel the old receipt PASS. The latest comparison explicitly normalizes only that one cache token while retaining the original full-65 hashes; primary/source and API before/after comparisons remain strict across all 65 tables.

The actual recovery 11 passed its declared owned backup/restore and API-only restart checks. The source-to-restored comparison normalizes only `retrieval_epochs.token`, retaining full-65 hashes and requiring exact equality for all 64 other tables and four original source files. Source before/after and API before/after restart comparisons retain strict full-65 equality. This is controlled cache-epoch invalidation, not unrestricted ignored differences; no host reboot, whole frontend/worker restart or another-machine recovery is established.

Review results

D09 recorded 23 physical requests, three native passes and one original same-route review-pass case across eight planned cases. Original PASS, FAIL, unavailable, disagreement and HTTP failure fields remain unchanged. Qwen G01's raw PASS retains a separate missing assistant-why explanation concern; no independent teaching certification follows. Two GLM HTTP 429 failures retain unknown-cost holds. Successful route metadata proves neither an independently attested endpoint variant nor meaningful answer quality.

New saved-tariff known-upper additions are CNY 0.126410 and USD 0.014061350. Actual protections remain DeepSeek CNY 82.901372 and OpenRouter USD 1.97. These are separate currencies and protections, not invoice-confirmed expenditure. Billing is null, unknown costs remain unavailable, no refund/hold release is inferred, and no new paid call is made by report preparation.

Next requirement-matched work

Review source regions for notation/table/reading-order meaning; finish current locator and revocation tests; verify compatible source reconstruction and dated exception statistics; include attributable concept/prerequisite review and the full-source dependency manifest in the final release.

Current release evidence

The common source is recorded in `PACKAGE_MANIFEST.json` and `Evidence/current-c9-freegate-summary.json`. Installed chronology is retained in `Evidence/installer-original-failed.json` and `continuation01-failed.json` through `continuation04-failed.json`, with the actual accepted continuation in `Evidence/continuation05-terminal.json`. Original failed recovery 10 is retained in `Evidence/original-recovery10-failed.json`; actual latest recovery is recorded in `Evidence/post-install-recovery-terminal.json`.

The complete archive and eight owner volumes are specified to use this same public operational payload. At content-writing time, package assembly, reconstruction, report copy equality and rendered-page QA were not yet verified. Acceptance of those later delivery results requires actual PACKAGE_MANIFEST.json and owner manifests for partitions, reconstruction and copy equality, together with Evidence/nine-reports-every-page-qa.json for inspection of every rendered page. These reports record the source and actual development outcomes at content-writing time; the linked external receipts establish later delivery verification when they exist and pass. This avoids circular DOCX and archive hashes.

Evidence and attribution

The public D09 eight-case summary is fingerprinted ``dc1d0790ee676799eab653f33f18acdb0dc2a504765f74938d48317b1efd81``. The failed d688 gate summary is ``54cab38d40419d56dfccfb61cc9266db44c0ed1cd1ea4b887d311a06ae73267e``; the earlier focused-6748 production compatibility bridge is ``e7a818a47ab0c82f42a842d0040105cb4f6f88bd03253b29106e647002d2817c``. The cumulative requirement trace is ``docs/execution/week09-current-requirements-trace-20261006.md``, SHA256 ``f387d4e5dadd426464208abe5851c66f3618839827c1696d25151d1088c40958``; the earlier source index remains dated provenance. The focused cleanup receipt is ``focus-execution-receipt-01.json``, SHA256 ``69b8c582bafb7c35e3354bfa6491b5ee24824665083b94a9d7f9edf68b095eba``, with JUnit SHA256 ``ce3f55056e96d22843af22f47796223a3e166e5b60c1ae863958e44af0029b52``. The formatting-only AST and Ruff receipt is ``ROOT-formatted-candidate-AST-and-ruff-receipt-02.json``, SHA256 ``206505861563cb356995b70dd44370c384649ac52431a47e1d2386d4e451c4b0``. The reports use aggregate receipts, without reading or exporting private test-fixture contents or individual metadata.

Accountable ownership is distinct from execution: the recorded implementation, tests, publication and model runs used the shared Codex workflow. No personal contribution, sign-off or human experiment is invented. Human reviews remain 0/null and paused; this limits research conclusions and does not block independent engineering. No current Canvas submission rule, course completion or learning gain is claimed.